

```
In [37]: import pandas as pd
data=pd.read_csv("/home/placement/Downloads/fiat500.csv")
```

```
In [38]: data.describe()
```

```
Out[38]:
```

|       | ID          | engine_power | age_in_days | km            | previous_owners | lat         | lon         | price        |
|-------|-------------|--------------|-------------|---------------|-----------------|-------------|-------------|--------------|
| count | 1538.000000 | 1538.000000  | 1538.000000 | 1538.000000   | 1538.000000     | 1538.000000 | 1538.000000 | 1538.000000  |
| mean  | 769.500000  | 51.904421    | 1650.980494 | 53396.011704  | 1.123537        | 43.541361   | 11.563428   | 8576.003901  |
| std   | 444.126671  | 3.988023     | 1289.522278 | 40046.830723  | 0.416423        | 2.133518    | 2.328190    | 1939.958641  |
| min   | 1.000000    | 51.000000    | 366.000000  | 1232.000000   | 1.000000        | 36.855839   | 7.245400    | 2500.000000  |
| 25%   | 385.250000  | 51.000000    | 670.000000  | 20006.250000  | 1.000000        | 41.802990   | 9.505090    | 7122.500000  |
| 50%   | 769.500000  | 51.000000    | 1035.000000 | 39031.000000  | 1.000000        | 44.394096   | 11.869260   | 9000.000000  |
| 75%   | 1153.750000 | 51.000000    | 2616.000000 | 79667.750000  | 1.000000        | 45.467960   | 12.769040   | 10000.000000 |
| max   | 1538.000000 | 77.000000    | 4658.000000 | 235000.000000 | 4.000000        | 46.795612   | 18.365520   | 11100.000000 |

```
In [39]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   ID              1538 non-null   int64
1   model           1538 non-null   object
2   engine_power    1538 non-null   int64
3   age_in_days     1538 non-null   int64
4   km              1538 non-null   int64
5   previous_owners 1538 non-null   int64
6   lat             1538 non-null   float64
7   lon             1538 non-null   float64
8   price           1538 non-null   int64
dtypes: float64(2), int64(6), object(1)
memory usage: 108.3+ KB
```

```
In [40]: data1=data.loc[(data.previous_owners==1)]
```

```
In [41]: data1
```

```
Out[41]:
```

|      | ID   | model  | engine_power | age_in_days | km     | previous_owners | lat       | lon       | price |
|------|------|--------|--------------|-------------|--------|-----------------|-----------|-----------|-------|
| 0    | 1    | lounge | 51           | 882         | 25000  | 1               | 44.907242 | 8.611560  | 8900  |
| 1    | 2    | pop    | 51           | 1186        | 32500  | 1               | 45.666359 | 12.241890 | 8800  |
| 2    | 3    | sport  | 74           | 4658        | 142228 | 1               | 45.503300 | 11.417840 | 4200  |
| 3    | 4    | lounge | 51           | 2739        | 160000 | 1               | 40.633171 | 17.634609 | 6000  |
| 4    | 5    | pop    | 73           | 3074        | 106880 | 1               | 41.903221 | 12.495650 | 5700  |
| ...  | ...  | ...    | ...          | ...         | ...    | ...             | ...       | ...       | ...   |
| 1533 | 1534 | sport  | 51           | 3712        | 115280 | 1               | 45.069679 | 7.704920  | 5200  |
| 1534 | 1535 | lounge | 74           | 3835        | 112000 | 1               | 45.845692 | 8.666870  | 4600  |
| 1535 | 1536 | pop    | 51           | 2223        | 60457  | 1               | 45.481541 | 9.413480  | 7500  |
| 1536 | 1537 | lounge | 51           | 2557        | 80750  | 1               | 45.000702 | 7.682270  | 5990  |
| 1537 | 1538 | pop    | 51           | 1766        | 54276  | 1               | 40.323410 | 17.568270 | 7900  |

1389 rows × 9 columns

```
In [42]: data1=data.drop(['ID','lat','lon'],axis=1)
```

```
In [43]: data1
```

```
Out[43]:
```

|      | model  | engine_power | age_in_days | km     | previous_owners | price |
|------|--------|--------------|-------------|--------|-----------------|-------|
| 0    | lounge | 51           | 882         | 25000  | 1               | 8900  |
| 1    | pop    | 51           | 1186        | 32500  | 1               | 8800  |
| 2    | sport  | 74           | 4658        | 142228 | 1               | 4200  |
| 3    | lounge | 51           | 2739        | 160000 | 1               | 6000  |
| 4    | pop    | 73           | 3074        | 106880 | 1               | 5700  |
| ...  | ...    | ...          | ...         | ...    | ...             | ...   |
| 1533 | sport  | 51           | 3712        | 115280 | 1               | 5200  |
| 1534 | lounge | 74           | 3835        | 112000 | 1               | 4600  |
| 1535 | pop    | 51           | 2223        | 60457  | 1               | 7500  |
| 1536 | lounge | 51           | 2557        | 80750  | 1               | 5990  |
| 1537 | pop    | 51           | 1766        | 54276  | 1               | 7900  |

1538 rows × 6 columns

```
In [44]: data1=pd.get_dummies(data)
```

In [45]: data1

Out[45]:

|      | ID   | engine_power | age_in_days | km     | previous_owners | lat       | lon       | price | model_lounge | model_pop | model_sport |
|------|------|--------------|-------------|--------|-----------------|-----------|-----------|-------|--------------|-----------|-------------|
| 0    | 1    | 51           | 882         | 25000  | 1               | 44.907242 | 8.611560  | 8900  | 1            | 0         | 0           |
| 1    | 2    | 51           | 1186        | 32500  | 1               | 45.666359 | 12.241890 | 8800  | 0            | 1         | 0           |
| 2    | 3    | 74           | 4658        | 142228 | 1               | 45.503300 | 11.417840 | 4200  | 0            | 0         | 1           |
| 3    | 4    | 51           | 2739        | 160000 | 1               | 40.633171 | 17.634609 | 6000  | 1            | 0         | 0           |
| 4    | 5    | 73           | 3074        | 106880 | 1               | 41.903221 | 12.495650 | 5700  | 0            | 1         | 0           |
| ...  | ...  | ...          | ...         | ...    | ...             | ...       | ...       | ...   | ...          | ...       | ...         |
| 1533 | 1534 | 51           | 3712        | 115280 | 1               | 45.069679 | 7.704920  | 5200  | 0            | 0         | 1           |
| 1534 | 1535 | 74           | 3835        | 112000 | 1               | 45.845692 | 8.666870  | 4600  | 1            | 0         | 0           |
| 1535 | 1536 | 51           | 2223        | 60457  | 1               | 45.481541 | 9.413480  | 7500  | 0            | 1         | 0           |
| 1536 | 1537 | 51           | 2557        | 80750  | 1               | 45.000702 | 7.682270  | 5990  | 1            | 0         | 0           |
| 1537 | 1538 | 51           | 1766        | 54276  | 1               | 40.323410 | 17.568270 | 7900  | 0            | 1         | 0           |

1538 rows × 11 columns

```
In [46]: y=data1['price']
x=data1.drop('price',axis=1)
```

In [47]:

y

```
Out[47]: 0      8900
          1      8800
          2      4200
          3      6000
          4      5700
          ...
        1533    5200
        1534    4600
        1535    7500
        1536    5990
        1537    7900
```

Name: price, Length: 1538, dtype: int64

```
In [48]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.33,random_state=42)
```

In [49]:

x\_test.head(5)

```
Out[49]:
```

|             | ID   | engine_power | age_in_days | km     | previous_owners | lat       | lon       | model_lounge | model_pop | model_sport |
|-------------|------|--------------|-------------|--------|-----------------|-----------|-----------|--------------|-----------|-------------|
| <b>481</b>  | 482  | 51           | 3197        | 120000 | 2               | 40.174702 | 18.167629 | 0            | 1         | 0           |
| <b>76</b>   | 77   | 62           | 2101        | 103000 | 1               | 45.797859 | 8.644440  | 0            | 1         | 0           |
| <b>1502</b> | 1503 | 51           | 670         | 32473  | 1               | 41.107880 | 14.208810 | 1            | 0         | 0           |
| <b>669</b>  | 670  | 51           | 913         | 29000  | 1               | 45.778591 | 8.946250  | 1            | 0         | 0           |
| <b>1409</b> | 1410 | 51           | 762         | 18800  | 1               | 45.538689 | 9.928310  | 1            | 0         | 0           |

```
In [50]: x_train.head(5)
```

```
Out[50]:
```

|     | ID  | engine_power | age_in_days | km    | previous_owners | lat       | lon      | model_lounge | model_pop | model_sport |
|-----|-----|--------------|-------------|-------|-----------------|-----------|----------|--------------|-----------|-------------|
| 527 | 528 | 51           | 425         | 13111 | 1               | 45.022388 | 7.58602  | 1            | 0         | 0           |
| 129 | 130 | 51           | 1127        | 21400 | 1               | 44.332531 | 7.54592  | 1            | 0         | 0           |
| 602 | 603 | 51           | 2039        | 57039 | 1               | 40.748241 | 14.52835 | 0            | 1         | 0           |
| 331 | 332 | 51           | 1155        | 40700 | 1               | 42.143860 | 12.54016 | 1            | 0         | 0           |
| 323 | 324 | 51           | 425         | 16783 | 1               | 41.903221 | 12.49565 | 1            | 0         | 0           |

```
In [51]: y_test.head(5)
```

```
Out[51]: 481      7900
76       7900
1502     9400
669      8500
1409     9700
Name: price, dtype: int64
```

```
In [52]: y_train.head(5)
```

```
Out[52]: 527      9990
129       9500
602       7590
331       8750
323       9100
Name: price, dtype: int64
```

```
In [53]: x_train.shape
```

```
Out[53]: (1030, 10)
```

```
In [54]: y_train
```

```
Out[54]: 527      9990
         129      9500
         602      7590
         331      8750
         323      9100
         ...
        1130     10990
        1294      9800
         860      5500
        1459      9990
        1126      8900
        Name: price, Length: 1030, dtype: int64
```

```
In [55]: #linear_regression
        from sklearn.linear_model import LinearRegression
        reg=LinearRegression()
        reg.fit(x_train,y_train)
```

```
Out[55]: LinearRegression()
In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.
```

```
In [56]: ypred=reg.predict(x_test)
```

In [57]: ypred

```
Out[57]: array([ 5819.19308764,  7248.82914161,  9741.8936974 ,  9798.98033074,
 10055.00624601,  9551.4955679 ,  9758.01743879, 10122.9778365 ,
  9654.9661814 ,  9251.1403257 , 10478.09512253,  7807.3005255 ,
  7705.15873781,  6295.63244894,  9545.40486313, 10422.92177704,
  9616.90811615,  7756.9171161 ,  4893.88454414, 10581.46142719,
 10465.24078346, 10443.29318231,  7518.43696046, 10028.21911459,
  6990.73118896,  8989.86900819,  4823.51364349,  6989.03118684,
  7822.83203734,  9683.17944083,  7344.21343132,  5341.43860798,
  5420.78405336,  5092.38401339,  8971.44357515,  5702.81242412,
  9920.16285466,  8334.58448277,  6220.93323723,  8389.23958511,
  9695.84208061,  6859.59630725,  9101.22635456, 10063.22592995,
  8621.83915759, 10175.06753933,  9063.21918346,  8867.24865352,
  7094.44228184,  9058.37693565,  9474.82390731, 10406.09102832,
 10112.65006224,  6820.90463865,  9700.36507783,  9382.18149429,
  9632.57617775, 10553.81356008,  9847.21129432,  7247.16814789,
  9990.23331336,  7084.23300123,  9977.34233656,  7245.01115798,
  6490.89305576,  9737.86785115,  9853.54349825,  8568.7125607 ,
  8506.81438703,  6484.69051659,  7883.1895563 ,  6870.28308427,
  8263.36833348, 10551.03496347,  7434.71134313,  8637.85174602,
  8762.87817027, 10010.47800077,  7224.68888828,  8527.72426022]
```



```
In [58]: Results=pd.DataFrame(columns=['price','predicted'])
Results['price']=y_test
Results['predicted']=ypred
Results=Results.reset_index()
Results['ID']=Results.index
Results.head(15)
```

```
Out[58]:
```

|    | index | price | predicted    | ID |
|----|-------|-------|--------------|----|
| 0  | 481   | 7900  | 5819.193088  | 0  |
| 1  | 76    | 7900  | 7248.829142  | 1  |
| 2  | 1502  | 9400  | 9741.893697  | 2  |
| 3  | 669   | 8500  | 9798.980331  | 3  |
| 4  | 1409  | 9700  | 10055.006246 | 4  |
| 5  | 1414  | 9900  | 9551.495568  | 5  |
| 6  | 1089  | 9900  | 9758.017439  | 6  |
| 7  | 1507  | 9950  | 10122.977837 | 7  |
| 8  | 970   | 10700 | 9654.966181  | 8  |
| 9  | 1198  | 8999  | 9251.140326  | 9  |
| 10 | 1088  | 9890  | 10478.095123 | 10 |
| 11 | 576   | 7990  | 7807.300526  | 11 |
| 12 | 965   | 7380  | 7705.158738  | 12 |
| 13 | 1488  | 6800  | 6295.632449  | 13 |
| 14 | 1432  | 8900  | 9545.404863  | 14 |

```
In [59]: Results['diff']=Results.apply(lambda row: row.price-row.predicted,axis=1)
```

In [60]: Results


Out[60]:

|     | index | price | predicted    | ID  | diff         |
|-----|-------|-------|--------------|-----|--------------|
| 0   | 481   | 7900  | 5819.193088  | 0   | 2080.806912  |
| 1   | 76    | 7900  | 7248.829142  | 1   | 651.170858   |
| 2   | 1502  | 9400  | 9741.893697  | 2   | -341.893697  |
| 3   | 669   | 8500  | 9798.980331  | 3   | -1298.980331 |
| 4   | 1409  | 9700  | 10055.006246 | 4   | -355.006246  |
| ... | ...   | ...   | ...          | ... | ...          |
| 503 | 291   | 10900 | 10121.593384 | 503 | 778.406616   |
| 504 | 596   | 5699  | 6288.648282  | 504 | -589.648282  |
| 505 | 1489  | 9500  | 10016.505537 | 505 | -516.505537  |
| 506 | 1436  | 6990  | 8248.746492  | 506 | -1258.746492 |
| 507 | 575   | 10900 | 10337.345820 | 507 | 562.654180   |

508 rows × 5 columns

```
In [61]: import seaborn as sns
import matplotlib.pyplot as plt
sns.lineplot(x='ID', y='price', data=Results.head(50))
sns.lineplot(x='ID', y='predicted', data=Results.head(50))
plt.plot()
```

Out[61]: []



```
In [62]: import warnings
warnings.filterwarnings('ignore')
```

```
In [63]: #ridge regression
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import Ridge
alpha = [1e-15, 1e-10, 1e-8, 1e-4, 1e-3, 1e-2, 1, 5, 10, 20, 30]
ridge = Ridge()
parameters = {'alpha': alpha}
ridge_regressor = GridSearchCV(ridge, parameters)
ridge_regressor.fit(x_train, y_train)
```

```
Out[63]: GridSearchCV(estimator=Ridge(),
                      param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.0001, 0.001, 0.01, 1,
                                             5, 10, 20, 30]})
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [64]: ridge_regressor.best_params_
```

```
Out[64]: {'alpha': 30}
```

```
In [65]: ridge=Ridge(alpha=30)
ridge.fit(x_train,y_train)
y_pred_ridge=ridge.predict(x_test)
```

```
In [66]: from sklearn.metrics import mean_squared_error
Ridge_Error=mean_squared_error(y_pred_ridge,y_test)
Ridge_Error
```

Out[66]: 574728.5696156605

```
In [67]: from sklearn.metrics import r2_score
r2_score(y_test,y_pred_ridge)
```

Out[67]: 0.8435021284061197

```
In [68]: Results=pd.DataFrame(columns=['Actual','Predicted'])
Results['Actual']=y_test
Results['Predicted']=y_pred_ridge
Results=Results.reset_index()
Results['ID']=Results.index
Results.head(10)
```


Out[68]:

|   | index | Actual | Predicted    | ID |
|---|-------|--------|--------------|----|
| 0 | 481   | 7900   | 5819.298540  | 0  |
| 1 | 76    | 7900   | 7264.574918  | 1  |
| 2 | 1502  | 9400   | 9738.882706  | 2  |
| 3 | 669   | 8500   | 9794.478395  | 3  |
| 4 | 1409  | 9700   | 10050.350724 | 4  |
| 5 | 1414  | 9900   | 9548.821263  | 5  |
| 6 | 1089  | 9900   | 9750.202837  | 6  |
| 7 | 1507  | 9950   | 10118.769447 | 7  |
| 8 | 970   | 10700  | 9656.236315  | 8  |
| 9 | 1198  | 8999   | 9247.205270  | 9  |

```
In [69]: import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [71]: sns.lineplot(x='ID',y='Actual',data=Results.head(50))  
sns.lineplot(x='ID',y='Predicted',data=Results.head(50))  
plt.plot()
```

Out[71]: []



```
In [72]: #elastic_model
from sklearn.linear_model import ElasticNet
from sklearn.model_selection import GridSearchCV
elastic = ElasticNet()
parameters = {'alpha': [1e-15, 1e-10, 1e-8, 1e-4, 1e-3, 1e-2, 1, 5, 10, 20]}
elastic_regressor = GridSearchCV(elastic, parameters)
elastic_regressor.fit(x_train, y_train)
```

```
Out[72]: GridSearchCV(estimator=ElasticNet(),
                      param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.0001, 0.001, 0.01, 1,
                                             5, 10, 20]})
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [73]: elastic_regressor.best_params_
```

```
Out[73]: {'alpha': 0.01}
```

```
In [74]: elastic=ElasticNet(alpha=.33)
elastic.fit(x_train,y_train)
y_pred_elastic=elastic.predict(x_test)
```

```
In [75]: from sklearn.metrics import r2_score
r2_score(y_test,y_pred_elastic)
```

```
Out[75]: 0.8445968963244241
```

```
In [76]: Results=pd.DataFrame(columns=['Actual','Predicted'])
Results['Actual']=y_test
Results['Predicted']=y_pred_elastic
Results=Results.reset_index()
Results['ID']=Results.index
Results.head(10)
```

```
Out[76]:
```


|   | index | Actual | Predicted    | ID |
|---|-------|--------|--------------|----|
| 0 | 481   | 7900   | 5834.887172  | 0  |
| 1 | 76    | 7900   | 7318.839756  | 1  |
| 2 | 1502  | 9400   | 9727.583531  | 2  |
| 3 | 669   | 8500   | 9778.566002  | 3  |
| 4 | 1409  | 9700   | 10033.013512 | 4  |
| 5 | 1414  | 9900   | 9538.968427  | 5  |
| 6 | 1089  | 9900   | 9721.786450  | 6  |
| 7 | 1507  | 9950   | 10102.881546 | 7  |
| 8 | 970   | 10700  | 9661.277720  | 8  |
| 9 | 1198  | 8999   | 9233.614930  | 9  |

```
In [77]: import seaborn as sns
import matplotlib.pyplot as plt
```



```
In [79]: sns.lineplot(x='ID',y='Actual',data=Results.head(50))  
sns.lineplot(x='ID',y='Predicted',data=Results.head(50))  
plt.plot()
```

Out[79]: []



In [ ]: